

# Hidden Markov Models (HMM)

**ECE/CS 498 DS U/G**

**Lecture 18: Hidden Markov Models continued**

Ravi K. Iyer

Dept. of Electrical and Computer Engineering

University of Illinois at Urbana Champaign

# Announcements

- Course Timeline
  - Mon 3/30: HMMs Continued, **ICA 4**
  - Wed 4/1: Intro to Factor Graphs
    - **Mid-semester Feedback** (to be completed via Google Form)
    - MP 3 released
- MP 2 Final Checkpoint due **tonight (March 30) @ 11:59 PM on Compass2G**
- Final Project
  - Progress report 2 due **Friday April 17 @ 11:59 PM** on Compass2G
    - There should be *substantial* progress with projects by this point (i.e. meaningful results, ML/AI models)

# Markov Model

- Consider a system which can occupy one of  $N$  discrete *states* or *categories*

$$x_t \in \{1, 2, \dots, N\} \longrightarrow \text{state at time } t$$

- We are interested in *stochastic* systems, in which state evolution is random
- Any *joint* distribution can be factored into a series of *conditional* distributions:

$$p(x_0, x_1, \dots, x_T) = p(x_0) \prod_{t=1}^T p(x_t \mid x_0, \dots, x_{t-1})$$

- For a *Markov* process, the next state depends only on the current state:

$$p(x_{t+1} \mid x_0, \dots, x_t) = p(x_{t+1} \mid x_t)$$

# State Transition Matrices

- A *stationary* Markov chain with  $N$  states is described by an  $N \times N$  *transition matrix*:


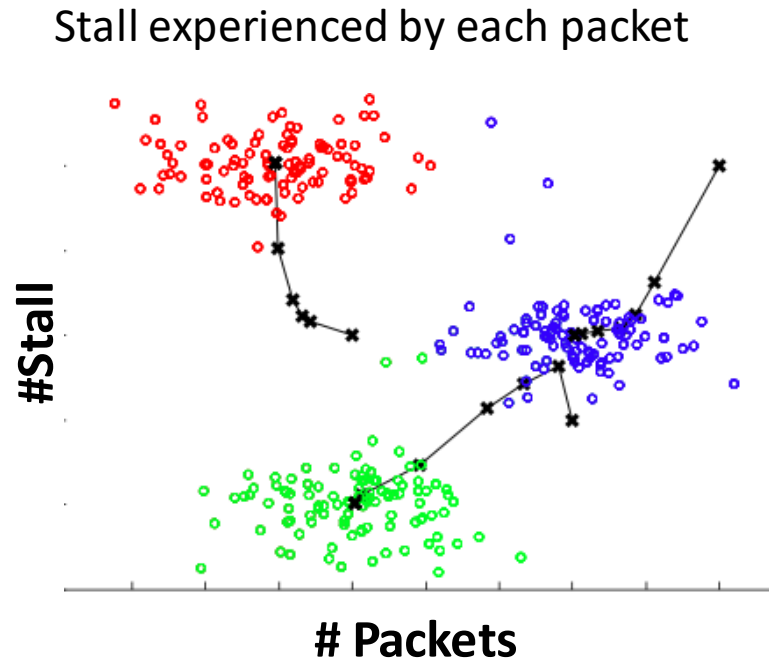
$$Q = \begin{bmatrix} q_{11} & q_{12} & q_{13} \\ q_{21} & q_{22} & q_{23} \\ q_{31} & q_{32} & q_{33} \end{bmatrix}$$

$$q_{ij} \triangleq p(x_{t+1} = i \mid x_t = j)$$

- Constraints on valid transition matrices:

$$q_{ij} \geq 0 \quad \sum_{i=1}^N q_{ij} = 1 \quad \text{for all } j$$

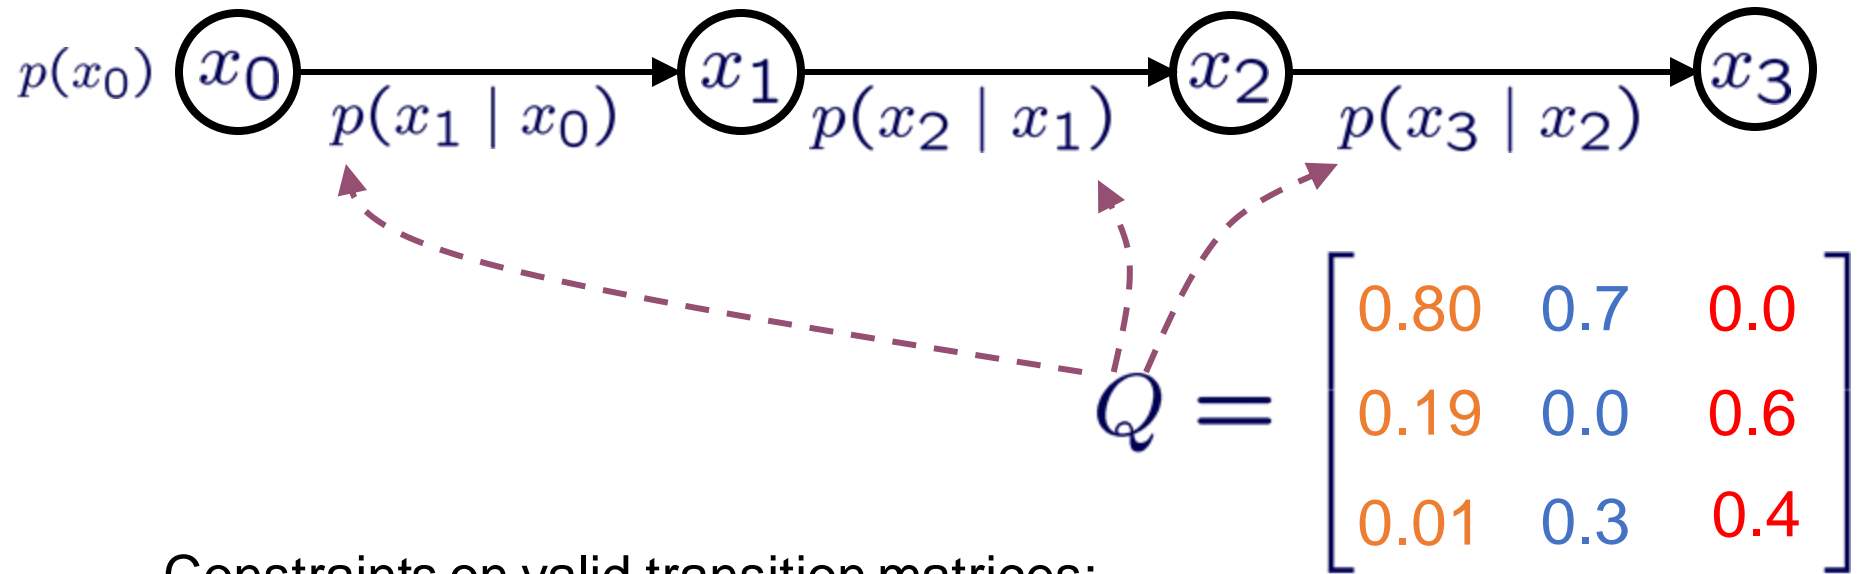
# State Transition Diagrams (Packet Stalls in a Network)



- Think of a particle randomly following an arrow at each discrete time step
- Most useful when  $N$  (# of States) small, and  $Q$  (TransProb Matrix) *sparse*

# Markov Chains: Graphical Models

$$p(x_0, x_1, \dots, x_T) = p(x_0) \prod_{t=1}^T p(x_t | x_{t-1})$$




Constraints on valid transition matrices:

$$q_{ij} \geq 0, \quad \sum_{i=1}^N q_{ij} = 1 \quad \text{for all } j$$

$$q_{ij} \triangleq p(x_{t+1} = i | x_t = j)$$

# Hidden Markov Models (Packet Stall Example Cont'd)

- Stall exists due to congestion
- Not directly measurable at runtime (hidden)
- Motivates *hidden Markov models* (HMM):

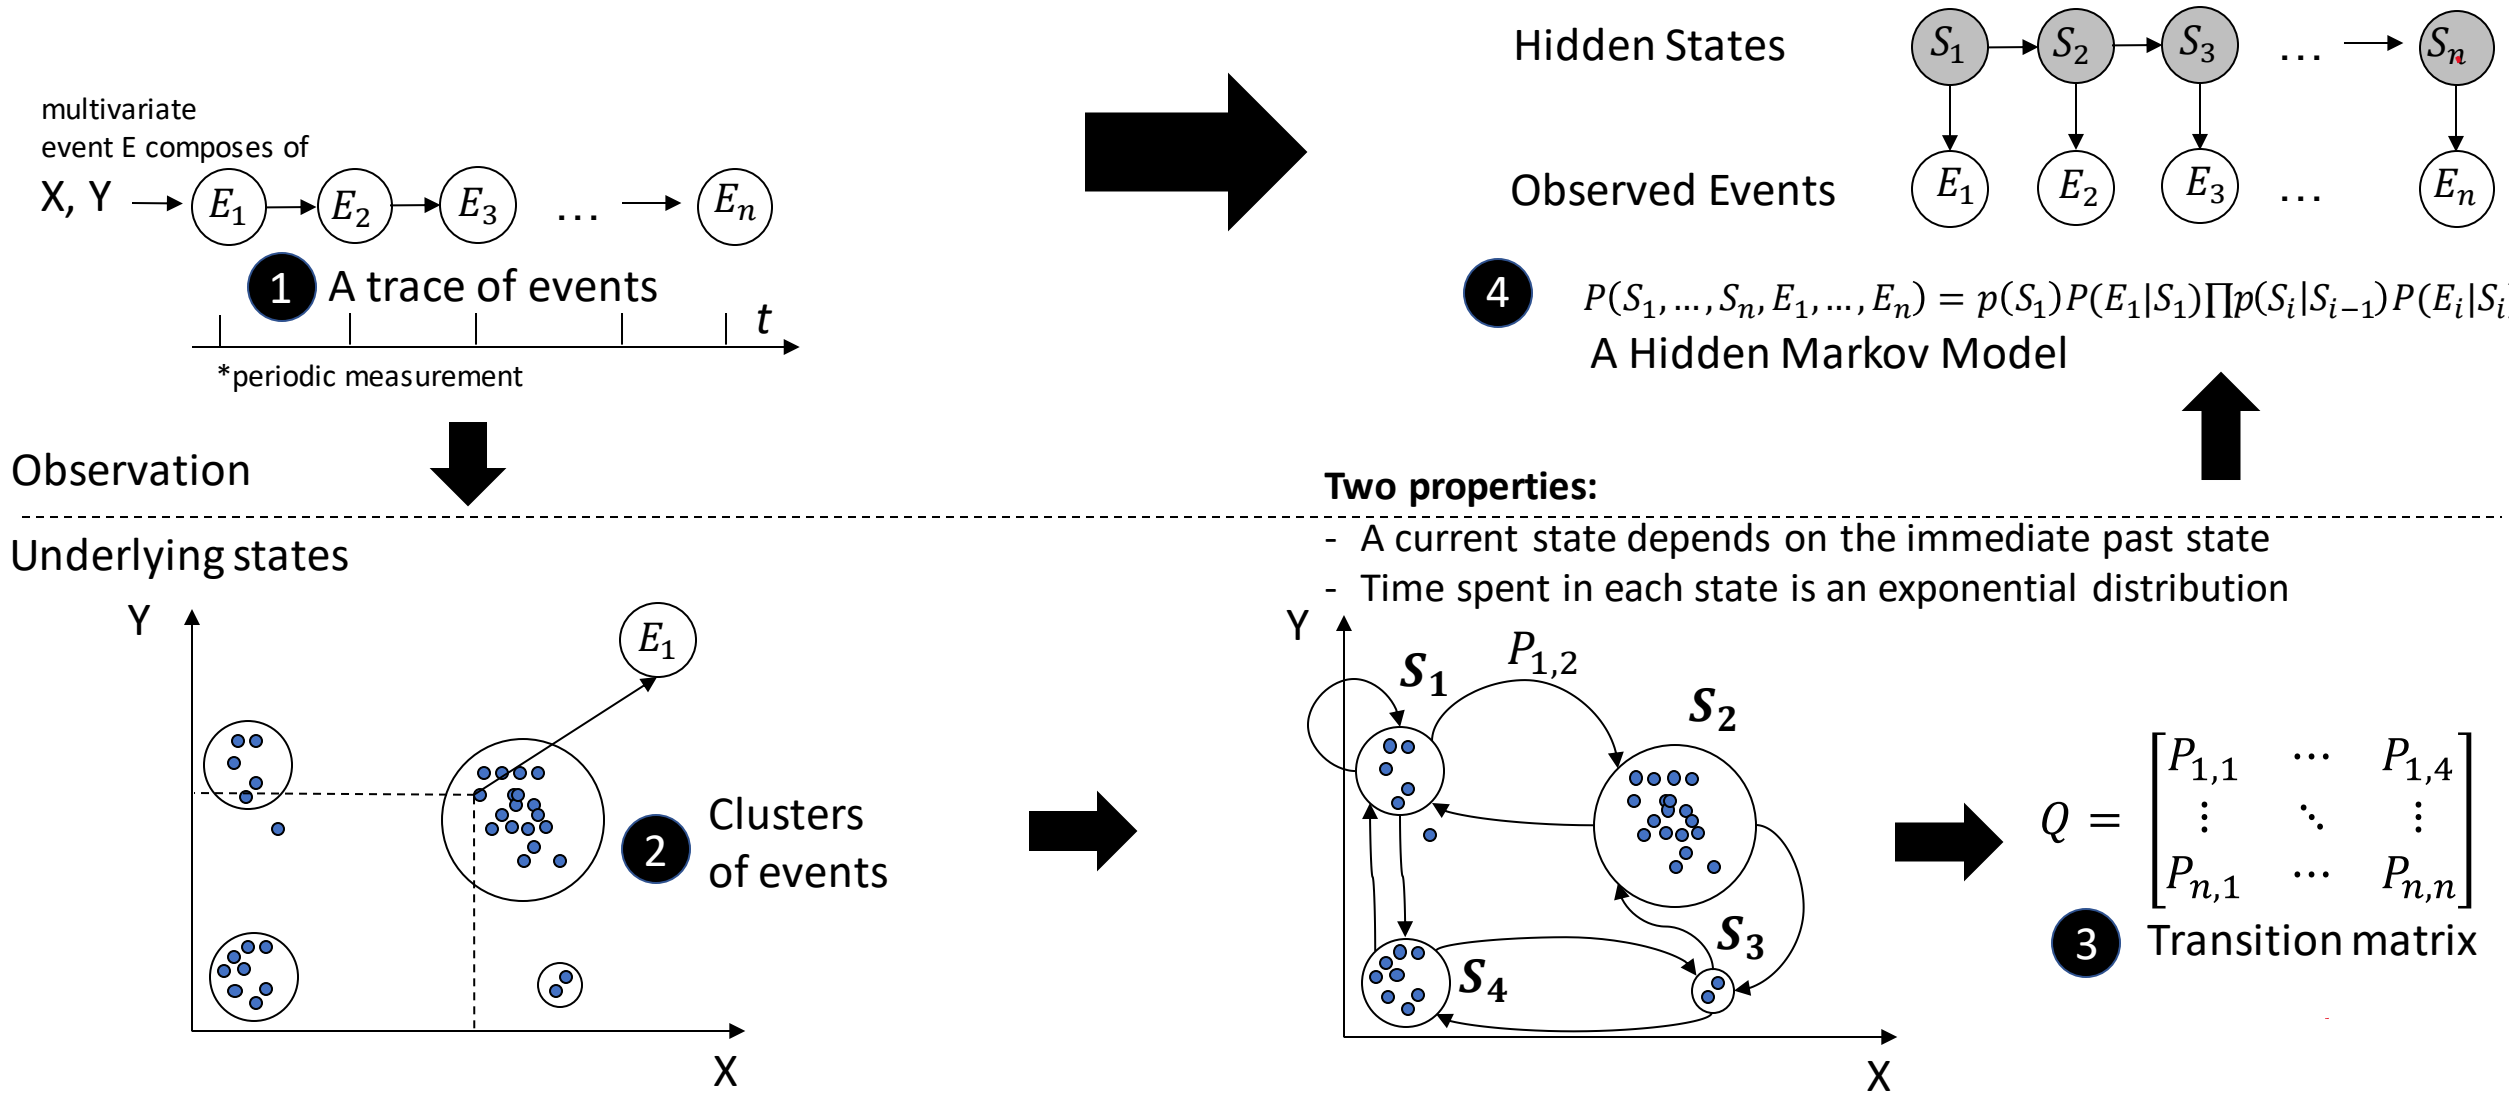


$$p(x_0, x_1, \dots, x_T) = p(x_0) \prod_{t=1}^T p(x_t | x_{t-1}) p(y_t | x_t)$$

Given  $x_t$ , previous observations impact future observations

$$p(y_t, y_{t+1}, \dots | x_t, y_{t-1}, y_{t-2}, \dots) = p(y_t, y_{t+1}, \dots | x_t)$$

# From a trace of events to a Hidden Markov Model





# Hidden Markov Models

## Model assumptions

An observation depends on its hidden state


A state variable only depends on the immediate previous state  
( Markov assumption)

The future observations and the past observations are  
conditionally independent given the current hidden state

## Advantages:

HMM can model sequential nature of input data (future  
depends on the past)

HMM has a linear-chain structure that clearly separates system  
state and observed events.




$$P(S_1, \dots, S_n, E_1, \dots, E_n) = p(S_1)P(E_1|S_1)\prod p(S_i|S_{i-1})P(E_i|S_i)$$

**A Hidden Markov model on observed events  
and system states**

# HMM Motivating Example: Paleontological Temperature Model

- Want to determine the average temperature at a particular place on earth over a sequence of years in the distant past
- Only annual average temperatures -- hot (**H**) and cold (**C**)
  - Probability of a hot year followed by another hot year is 0.7, and the probability of a cold year followed by another cold year is 0.6, independent of the temperature in prior years
- Correlation between the size of tree growth rings and temperature
  - Three different ring sizes, small (**T**), medium (**D**), and large (**L**)
- Assume that probability values from current period held in paleontological period too
- Determine the most likely temperature state in past years
  - Can't directly observe the temperature in the past
  - We can observe the size of tree rings – can this information be used?



H

C



T

D

L


Tree ring size

# Paleontological Temperature Model

- State space of hidden states:  $S = \{H, C\}$
- State space of observations:  $E = \{T, D, L\}$


- Transition probability matrix:  $A$
- Observation Matrix:  $B$
- Initial distribution for the hidden states:  $\pi$

Given by an oracle



# Paleontological Temperature Model

Example sequence with 5 observations



Determine the sequence of hidden states

# Hidden Markov Models

## Model


- Set of hidden states  $\mathcal{S} = \{\sigma_1, \dots, \sigma_N\}$
- Set of observable events  $\mathcal{E} = \{\epsilon_1, \dots, \epsilon_M\}$
- Transition probability matrix  $A$
- Observation matrix  $B$
- Initial distribution of hidden states  $\pi$

## Model assumptions

- An observation depends on its hidden state
- A state variable only depends on the immediate previous state (Markov assumption)
- The future observations and the past observations are **conditionally independent** given the current hidden state

## Advantages:

- HMM can model sequential nature of input data (future depends on the past)
- HMM has a linear-chain structure that clearly separates system state and observed events.



**A Hidden Markov model on observed events and system states**

$$\begin{aligned} &P(S_1, \dots, S_n, E_1, \dots, E_n) \\ &= P(S_1)P(E_1|S_1) \prod_{i=2}^n P(S_i|S_{i-1})P(E_i|S_i) \end{aligned}$$

# Inference question – Paleontological Temperature

Given the sequence of 5 observations  $T, L, D, T, D$  and the model  $(A, B, \pi)$ , how do we choose a corresponding state sequence  $S_1, S_2, \dots, S_n$  which is optimal in some meaningful sense (i.e., best explains the observations) where  $S_t \in \{H, C\}$ ?

**A simpler question:** Given the sequence of 5 observations  $T, L, D, T, D$  and the model  $(A, B, \pi)$ , which of the two is more probable eg.,  $S_3 = H$  or  $S_3 = C$ ?

# HMM Security Example

- Suppose you are a security expert monitoring the NCSA system
- By monitoring the system events, you want to say whether the system is safe or not
  - System's safety is a hidden state
  - Events are observed
  - Events are related to the safety of the system
- Is the system safe?
  - **HMM** to the rescue!


# Security Example: Transition Matrix

## Transition matrix (A)

The system has three distinct security states –

- (a) No Attack (**NA**),
- (b) Attack in Progress (**AP**), and
- (c) Attack Complete (**AC**).

- Every hour, the system is being attacked by attackers coordinating together around the world and trying to compromise the system.
- The system states always transition from **NA to AP** and **AP to AC**.
- An attacker is successful in changing the state of the system with probability of 0.75 and fails with a probability of 0.25.
- If the attack fails, the system stays in its current state.
- If the system state reaches **AC** the attack is complete, and the system stays in that state.



$$A = \begin{matrix} & \begin{matrix} \text{NA} & \text{AP} & \text{AC} \end{matrix} \\ \begin{matrix} \text{NA} \\ \text{AP} \\ \text{AC} \end{matrix} & \begin{pmatrix} 0.25 & 0.75 & 0 \\ 0 & 0.25 & 0.75 \\ 0 & 0 & 1 \end{pmatrix} \end{matrix}$$

Transition Probability Matrix



# Security Example: Emission matrix and initial distribution

## Observation matrix (**B**)

- Your monitoring system reports two types of events
  - Port Scan (**PS**)
  - Software Installation (**SI**)
- Monitors are always accurate and works. Attackers cannot compromise the monitors. Every hour, we get information from the monitors if the attackers are trying to do **PS or SI**.

## Initial distribution ( $\pi$ )

- We have no idea about the initial state of the system.

$$\mathbf{B} = \begin{array}{c} \mathbf{NA} \\ \mathbf{AP} \\ \mathbf{AC} \end{array} \begin{array}{cc} \mathbf{PS} & \mathbf{PI} \\ \left( \begin{array}{cc} P_{PS|NA} & P_{PI|NA} \\ P_{PS|AP} & P_{PI|AP} \\ P_{PS|AC} & P_{PI|AC} \end{array} \right) \end{array}$$

Observation Matrix

$$\pi_0 = \begin{array}{c} \mathbf{NA} \\ \mathbf{AP} \\ \mathbf{AC} \end{array} \left( \begin{array}{c} 1/3 \\ 1/3 \\ 1/3 \end{array} \right)$$

Initial state  
distribution/prior

# General Inference question

Given the sequence of  $n$  observations  $E_1, E_2, \dots, E_n$ , and the model  $(A, B, \pi)$ , how do we choose a corresponding state sequence  $S_1, S_2, \dots, S_n$  which is optimal in some meaningful sense (i.e., best explains the observations)?

**A simpler question:** Given the sequence of  $n$  observations  $E_1, E_2, \dots, E_n$ , and the model  $(A, B, \pi)$ , what is the most probable state  $S_t$  at  $t \in \{1, \dots, n\}$ ?

$$\operatorname{argmax}_{j \in \{1, \dots, N\}} P(S_t = \sigma_j | E_1, E_2, \dots, E_n)$$

$$S = \{\sigma_1, \dots, \sigma_N\}$$

# Breaking down the inference question

$$\begin{aligned} P(S_t | E_1, E_2, \dots, E_n) &= \frac{P(S_t, E_1, \dots, E_n)}{P(E_1, \dots, E_n)} = \frac{P(S_t, E_1, \dots, E_t, E_{t+1}, \dots, E_n)}{P(E_1, \dots, E_n)} \\ &= \frac{P(E_{t+1}, \dots, E_n | S_t, E_1, \dots, E_t) P(S_t, E_1, \dots, E_t)}{P(E_1, \dots, E_n)} \\ &= P(E_{t+1}, \dots, E_n | S_t, E_1, \dots, E_t) P(S_t | E_1, \dots, E_t) \frac{P(E_1, \dots, E_t)}{P(E_1, \dots, E_n)} \\ &= \frac{P(E_{t+1}, \dots, E_n | S_t) P(S_t | E_1, \dots, E_t)}{P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)} \end{aligned}$$

Bayes rule

Bayes rule

Markov property

# Breaking down the inference question

$$P(S_t | E_1, E_2, \dots, E_n) = \frac{P(E_{t+1}, \dots, E_n | S_t) P(S_t | E_1, \dots, E_t)}{P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)}$$

$P(S_t | E_1, \dots, E_t)$ :

Probability of hidden state at time  $t$  given observation up to time  $t$  (**Forwards algorithm**)

$P(E_{t+1}, \dots, E_n | S_t)$ :

Probability of the future observed sequence given the hidden state at time  $t$  (**Backwards algorithm**)


$P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)$ :

Does not depend on the hidden state (will not affect the maximization because it is just a scaling factor)

# Forwards algorithm: Paleontological Temperature

Want to calculate  $P(S_t|E_1, \dots, E_t)$

- Let us calculate it for  $t = 2$
- In the example,  $E_1 = T, E_2 = L$
- Find  $P(S_2 = H|E_1 = T, E_2 = L)$ ?



$$P(S_2 = H|E_1 = T, E_2 = L) = \frac{P(S_2 = H, E_1 = T, E_2 = L)}{P(E_1 = T, E_2 = L)}$$

Adding hidden state  $S_1$

$$= \frac{\sum_{s \in \{H, C\}} P(S_2 = H, E_1 = T, E_2 = L, S_1 = s)}{P(E_1 = T, E_2 = L)}$$

# Forwards algorithm: Paleontological Temperature

$$\begin{aligned}
 & \frac{\sum_{s \in \{H, C\}} P(\mathbf{S}_2 = H, E_1 = T, E_2 = L, \mathbf{S}_1 = s)}{P(E_1 = T, E_2 = L)} \\
 &= \frac{\sum_{s \in \{H, C\}} P(E_2 = L | \mathbf{S}_2 = H, E_1 = T, \mathbf{S}_1 = s) P(\mathbf{S}_2 = H, E_1 = T, \mathbf{S}_1 = s)}{P(E_1 = T, E_2 = L)} \quad \text{Bayes rule} \\
 & \quad \text{Markov property} \downarrow \\
 &= \frac{\sum_{s \in \{H, C\}} P(E_2 = L | \mathbf{S}_2 = H) P(\mathbf{S}_2 = H | E_1 = T, \mathbf{S}_1 = s) P(\mathbf{S}_1 = s | E_1 = T) P(E_1 = T)}{P(E_1 = T, E_2 = L)} \quad \text{Bayes rule} \\
 & \quad \text{Markov property} \downarrow \\
 &= \frac{\sum_{s \in \{H, C\}} P(E_2 = L | \mathbf{S}_2 = H) P(\mathbf{S}_2 = H | \mathbf{S}_1 = s) P(\mathbf{S}_1 = s | E_1 = T)}{P(E_2 = L | E_1 = T)} \quad \text{Bayes rule}
 \end{aligned}$$

# Forwards algorithm: Paleontological Temperature

Hidden state given all observations up to that point

Observation probability

Transition probability

Hidden state given all observations up to that point

$$P(S_2 = H | E_1 = T, E_2 = L) = \frac{P(E_2 = L | S_2 = H) \sum_{s \in \{H, C\}} P(S_2 = H | S_1 = s) P(S_1 = s | E_1 = T)}{P(E_2 = L | E_1 = T)}$$

Define:  $\alpha_t(i) = P(S_t = \sigma_i | E_1, E_2, \dots, E_t)$  and  $Z_t = P(E_t | E_1, \dots, E_{t-1})$

Above equation can be written as,

$$\alpha_2(H) = \frac{1}{Z_2} P(E_2 = L | S_2 = H) \sum_{s \in \{H, C\}} P(S_2 = H | S_1 = s) \alpha_1(s)$$

Where,  $Z_2 = P(E_2 | E_1)$

**Recursion**

# Forwards algorithm: General Expression

Define:  $\alpha_t(j) = P(S_t = \sigma_j | E_1, E_2, \dots, E_t)$  and  $Z_t = P(E_t | E_1, \dots, E_{t-1})$

In general,

$$\alpha_t(j) = \frac{1}{Z_t} P(E_t | S_t = \sigma_j) \sum_{i=1}^N P(S_t = \sigma_j | S_{t-1} = \sigma_i) \alpha_{t-1}(i) \quad Z_t = \sum_{j=1}^N b_t \odot (A^T \alpha_{t-1})$$

Transition probability  $a_{ij}$

Above equation can be written as a matrix for all  $j$ ,

$$\begin{bmatrix} \alpha_t(1) \\ \vdots \\ \alpha_t(j) \\ \vdots \\ \alpha_t(N) \end{bmatrix} \propto \begin{bmatrix} P(E_t | S_t = \sigma_1) \\ \vdots \\ P(E_t | S_t = \sigma_j) \\ \vdots \\ P(E_t | S_t = \sigma_N) \end{bmatrix} \odot \begin{bmatrix} a_{11} & \dots & \dots & \dots & a_{N1} \\ \vdots & \ddots & \dots & \dots & \vdots \\ a_{1j} & \dots & a_{ij} & \dots & a_{Nj} \\ \vdots & \dots & \dots & \ddots & \dots \\ a_{1N} & \dots & \dots & \dots & a_{NN} \end{bmatrix} \begin{bmatrix} \alpha_{t-1}(1) \\ \vdots \\ \alpha_{t-1}(i) \\ \vdots \\ \alpha_{t-1}(N) \end{bmatrix}$$

⊙ Represents elementwise product (Hadamard product)

$$\alpha_t \propto b_t \odot (A^T \alpha_{t-1})$$

$b_t$  is the column of the observation matrix B corresponding to  $E_t$



# Forwards Algorithm: Paleontological Temperature

For observations  $T, L, D, T, L$

$P(S_2 | E_1 = T, E_2 = L)$  is,

$$\begin{bmatrix} \alpha_2(H) \\ \alpha_2(C) \end{bmatrix} \propto \begin{bmatrix} 0.5 \\ 0.1 \end{bmatrix} \odot \left( \begin{bmatrix} 0.7 & 0.4 \\ 0.3 & 0.6 \end{bmatrix} \begin{bmatrix} \alpha_1(H) \\ \alpha_1(C) \end{bmatrix} \right)$$

	H	C
H	0.7	0.3
C	0.4	0.6

Transition probability matrix

	T	D	L
H	0.1	0.4	0.5
C	0.7	0.2	0.1

Observation matrix

Similarly,  $P(S_3 | E_1 = T, E_2 = L, E_3 = D)$  is,

$$\begin{bmatrix} \alpha_3(H) \\ \alpha_3(C) \end{bmatrix} \propto \begin{bmatrix} 0.4 \\ 0.2 \end{bmatrix} \odot \left( \begin{bmatrix} 0.7 & 0.4 \\ 0.3 & 0.6 \end{bmatrix} \begin{bmatrix} \alpha_2(H) \\ \alpha_2(C) \end{bmatrix} \right)$$

# Forwards Algorithm

1. Input:  $(A, B, \pi)$  and observed sequence  $E_1, \dots, E_n$
2.  $[\alpha_1, Z_1] = \text{normalize}(b_1 \odot \pi)$
3. **for**  $t = 2:n$  **do**  
     $[\alpha_t, Z_t] = \text{normalize}(b_t \odot (A^T \alpha_{t-1}))$
4. return  $\alpha_1, \dots, \alpha_n$  and  $\log(P(E_1, \dots, E_n)) = \sum_t \log(Z_t)$

Note:

Subroutine:  $[v, Z] = \text{normalize}(u)$ :  $Z = \sum_j u_j$ ;  $v_j = u_j/Z$ ;

# Breaking down the inference question

$$P(S_t | E_1, E_2, \dots, E_n) = \frac{P(E_{t+1}, \dots, E_n | S_t) P(S_t | E_1, \dots, E_t)}{P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)}$$

$P(S_t | E_1, \dots, E_t)$ :

Probability of hidden state at time  $t$  given observation up to time  $t$  (**Forwards algorithm**)

$P(E_{t+1}, \dots, E_n | S_t)$ :

Probability of the future observed sequence given the hidden state at time  $t$  (**Backwards algorithm**)

$P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)$ :

Does not depend on the hidden state (will not affect the maximization because it is just a scaling factor)

# Backwards Algorithm (similar to Forwards Algo.)

Calculate  $P(E_{t+1}, \dots, E_n | S_t)$

Define:  $\beta_t(j) = P(E_{t+1}, \dots, E_n | S_t = \sigma_j)$

Include  $S_t$  to use information from the one-step future

$$\begin{aligned}
 \beta_{t-1}(j) &= P(E_t, \dots, E_n | S_{t-1} = \sigma_j) = \sum_{i=1}^N P(S_t = \sigma_i, E_t, \dots, E_n | S_{t-1} = \sigma_j) \\
 &= \sum_{i=1}^N P(E_{t+1}, \dots, E_n | S_{t-1} = \sigma_j, S_t = \sigma_i, E_t) P(E_t | S_{t-1} = \sigma_j, S_t = \sigma_i) P(S_t = \sigma_i | S_{t-1} = \sigma_j) \\
 &= \sum_{i=1}^N P(E_{t+1}, \dots, E_n | S_t = \sigma_i) P(E_t | S_t = \sigma_i) P(S_t = \sigma_i | S_{t-1} = \sigma_j) \\
 &= \sum_{i=1}^N \beta_t(i) P(E_t | S_t = \sigma_i) P(S_t = \sigma_i | S_{t-1} = \sigma_j)
 \end{aligned}$$

Chain rule

Markov property

By definition of  $\beta_t(j)$

Emission probability

Transition probability

In matrix form, we get,

$$\beta_{t-1} = A(b_t \odot \beta_t)$$

$$\beta_t = \begin{bmatrix} \beta_t(1) \\ \vdots \\ \beta_t(N) \end{bmatrix}$$

# Backwards Algorithm

1. Input:  $(A, B, \pi)$  and observed sequence  $E_1, \dots, E_n$
2.  $\beta_n = 1$  ; // initialize  $\beta_n(j)$  to 1 for all states  $\sigma_j$
3. **for**  $t = n - 1 : 1$  **do**  
     $\beta_{t-1} = A(b_t \odot \beta_t)$
4. return  $\beta_1, \dots, \beta_n$

# Breaking down the inference question

$$P(S_t | E_1, E_2, \dots, E_n) = \frac{P(E_{t+1}, \dots, E_n | S_t) P(S_t | E_1, \dots, E_t)}{P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)}$$

$P(S_t | E_1, \dots, E_t)$ :

Probability of hidden state at time  $t$  given observation up to time  $t$  (**Forwards algorithm**)

$P(E_{t+1}, \dots, E_n | S_t)$ :

Probability of the future observed sequence given the hidden state at time  $t$  (**Backwards algorithm**)

$P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)$ :

Does not depend on the hidden state (will not affect the maximization because it is just a scaling factor)

# Inference – using Forwards-Backwards expressions

$$P(S_t | E_1, E_2, \dots, E_n) = \frac{P(E_{t+1}, \dots, E_n | S_t) P(S_t | E_1, \dots, E_t)}{P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)}$$

For  $S_t = \sigma_j$  and  $\gamma_t(j) = P(S_t = \sigma_j | E_1, E_2, \dots, E_n)$ , the above equation is:

$$P(S_t = \sigma_j | E_1, E_2, \dots, E_n) = \frac{P(E_{t+1}, \dots, E_n | S_t = \sigma_j) P(S_t = \sigma_j | E_1, \dots, E_t)}{P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)}$$

$$\gamma_t(j) = \frac{\beta_t(j) \alpha_t(j)}{P(E_{t+1}, \dots, E_n | E_1, \dots, E_t)} = \frac{\beta_t(j) \alpha_t(j)}{\sum_{i=1}^N \beta_t(j) \alpha_t(j)}$$



Theorem of total probability

$\gamma_t(j) \propto \beta_t(j) \alpha_t(j)$

# Inference: Most likely state

- Forwards-backwards algorithm gives  $P(S_t = \sigma_j | E_1, \dots, E_n)$  for all  $j$
- Find the **individually most likely state** at time  $t$  given all observations

$$S_t^* = \operatorname{argmax}_{j \in \{1, \dots, N\}} \gamma_t(j)$$



# Optimality of inference

- In the inference problem we attempt to uncover the hidden part of HMM, i.e., find the “correct” state sequence
- It is impossible to find the “correct” state sequence (solution)
- Use optimality criterion to find the “best” possible solution
- **Several reasonable criteria** exist and is a strong function of the intended application
  - **Most likely state given observations**
    - Application in finding average statistics, expected number of correct states
    - Solved using **Forwards-Backwards algorithm**
  - **Single best sequence that maximises probability of observed events**
    - Application in continuous speech recognition
    - Solved using **Viterbi algorithm**

# Resources

Rabiner's (excellent) paper:

<https://www.ece.ucsb.edu/Faculty/Rabiner/ece259/Reprints/tutorial%20on%20hmm%20and%20applications.pdf>