

# Cross Validation, Reinforcement Learning

## Lecture 27: Intro to Reinforcement Learning Cross Validation

ECE/CS 498 DS

Professor Ravi K. Iyer

Department of Electrical and Computer  
Engineering

University of Illinois

# Announcements

- MP 2 final grades and MP 3 checkpoint 1 grades have been released
- MP 3 final submission due **tonight 4/29 @ 11:59 PM** via Compass
- HW 5 released, due **Mon 5/4 @ 11:59 PM** via Compass
  - Covers SVM, neural networks, and random forest
- Discussion section on **Fri 5/1** will be additional practice with neural networks



# Final Exam Logistics

# Final Exam Logistics

- Final Exam will **be Friday 5/15 from 8-11:10 AM**
- Exam length:
  - You will have 10 minutes of dedicated reading time from 8-8:10 AM
  - After that, exam is designed to take ~2.5 hours (you have 3 hours)
  - You need to scan and submit your solutions by 11:10 AM on Compass
- You are allowed 3 physical, double-sided, 8.5x11 cheat sheets (can be typed or handwritten)
- **No calculators** or other electronic devices are allowed
- **No reference** materials/books are allowed

# Final Exam Logistics

- **Content:**
  - **All course material is fair game**
  - **Material from after midterm exam will be emphasized:**
    - HMMs and Forward-Backward Algorithm
    - Factor Graphs and Belief Propagation
    - SVM
    - Neural Networks
    - Random Forest
    - Cross Validation, etc.
  - **Earlier course material may be included:**
    - Basic Probability
    - Naïve Bayes and Bayesian Networks
    - Clustering (K-means, GMM, Hierarchical)
    - Linear/Logistic Regression
    - PCA, etc.

# Final Exam Logistics

- Exam will be released as PDF via Compass, and you will have until 11:10 AM to (i) complete it and (ii) scan and submit your solutions on Compass
  - **Submission format** – preferably .pdf, but .jpg/.jpeg will be accepted as well.
  - **Submission quality** – Make sure to use the highest-quality files available to you. You will be responsible for ensuring that all scanned files are clearly legible/readable by course staff. We cannot grade what we cannot read...
  - **Strict deadline** – You must scan and submit your solutions by 11:10 AM
  - **Smartphone Scanning** – If needed, you may use your smartphone to scan your exam after you have finished writing
    - Writing after using smartphone will be considered cheating, and you are not allowed to use your smartphone at any other point during the exam
    - Recommended smartphone apps for scanning: Genius Scan, CamScanner

# Final Exam Logistics

- Exam will be proctored remotely using **Proctorio**.  
Requirements:
  - Install **Google Chrome**
  - Install **Proctorio extension** for Chrome (<https://getproctorio.com/>)
  - Have your **iCard** (or some other official ID) available to show the camera at the beginning of the exam
  - Have **~20 blank sheets** of white printer paper to write on during the exam. You may not need all 20, but it's better to have extra
  - Have a **webcam** available during the exam (laptop webcam is fine)
    - All students must have a webcam to use during the exam for us to verify academic integrity
  - Have a **microphone** available during the exam (laptop microphone is fine)
- Webcam should be setup to show you working at a desk/table



# Cross Validation



# Validation:

## Binary Classification Performance

You are solving a Binary Classification Problem, Predict if a person has cancer (1) or not (0). You have learned a model using your training dataset.

**How do you test your model?**

Explanation:

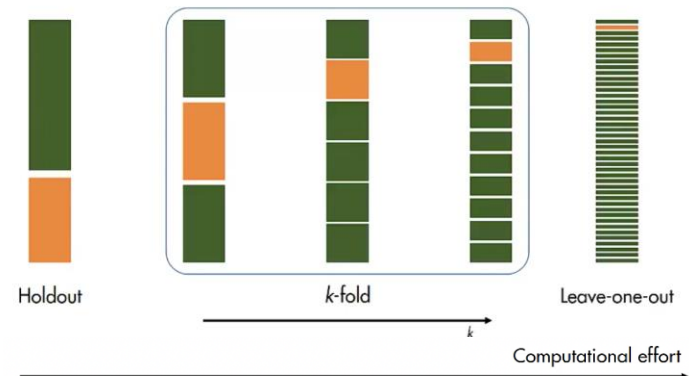
- 1) **Holdout:** Randomly divide dataset into 2 parts, learn model on training data and validate on the test data.
- 2) **K-fold:** Divide data into K equal parts, run the following K-times: (hold out one of the K parts of the dataset, train on the remaining, test on the holdout). Every data-sample is part of the training data K-1 times and part of the test data 1 time.
- 3) **Leave-one-out:** For every data sample, leave it out, train the model on the remaining data and predict for the left-out data sample.
- 4) **Monte Carlo Cross Validation:** In each of multiple trials, randomly divide the dataset into training and testing sets (of prespecified sizes). Always retrain the model on the trial's training set and validate on the trial's testing set.

**Cross Validation:** For a given training iteration, divide the dataset into 2 parts:

- training set (for training the model)
- test set (for validating the model)

Different ways of doing this:

- (1) Holdout
- (2) K-fold
- (3) Leave-one-out
- (4) Monte Carlo



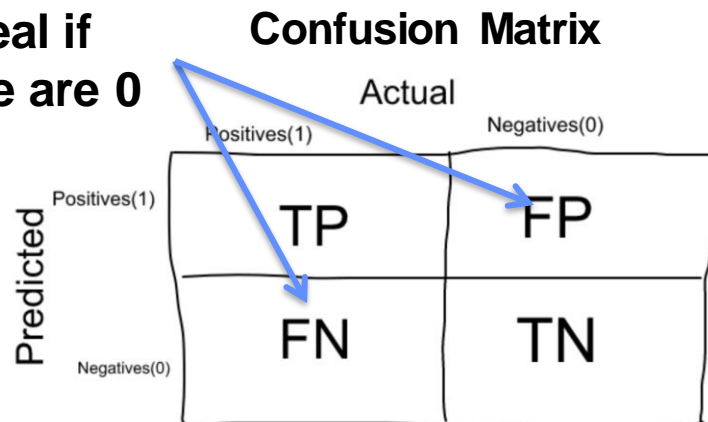
**green – training data**  
**orange - test data**

# What does it mean to **validate** on the test data?

Ideal if  
these are 0

**Confusion Matrix**

		Actual	
		Positives(1)	Negatives(0)
Predicted	Positives(1)	TP	FP
	Negatives(0)	FN	TN



Depending on the **cost of making a mistake**, minimize the number of FN or FP

## Common issues

**Precision** = 1 if FP=0, but FN may be high

**Sensitivity** = 1 if predict everything as positive

**Specificity** = 1 if predict everything as negative

How to strike a **balance**? Use F1-score!

## Metrics

$$\text{Accuracy} = \frac{TP + TN}{TP + FP + TN + FN}$$

$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Sensitivity} = \frac{TP}{TP + FN}$$

(Recall)

$$\text{Specificity} = \frac{TN}{TN + FP}$$

$$\text{F1-score} = \frac{2 * \text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

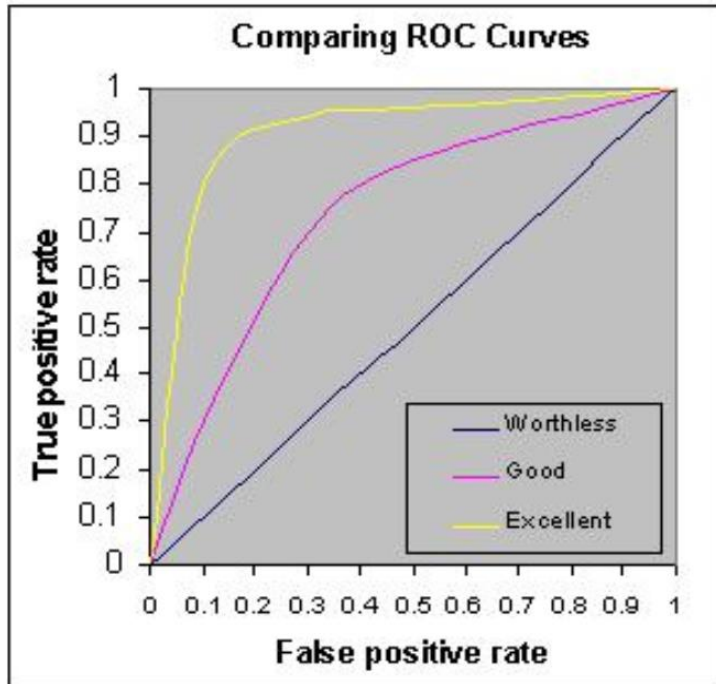
**Harmonic mean** of Precision and Recall  
(when large disparity between the 2 values,  
the score is closer to the smaller value)

# Why not Always Use Accuracy?

- Suppose you need to perform binary classification, where the classes are either positive or negative
- If the class distribution is imbalanced, another metric called AUC may be a better performance evaluation criterion
  - Let's say 90% of the test labels are positive
  - A very simple model that only predicts positive labels (and doesn't use the training data at all) would achieve 90% accuracy
    - If the training data is also just as skewed, then even a more complex model may end up predicting positive most of the time
  - This high accuracy is misleading since it depends solely on the distribution of the test dataset
  - By capturing results after varying the decision threshold  $\tau$  in the range  $[0, 1]$ , AUC measures discriminative performance of the model
    - Thus, AUC is generally preferred for evaluating binary classifiers with imbalanced datasets
    - The very simple model would report a poor AUC score

# ROC and AUC

Receiver Operating Characteristic curve: A plot of **False Positive Rate** Vs **True Positive Rate**

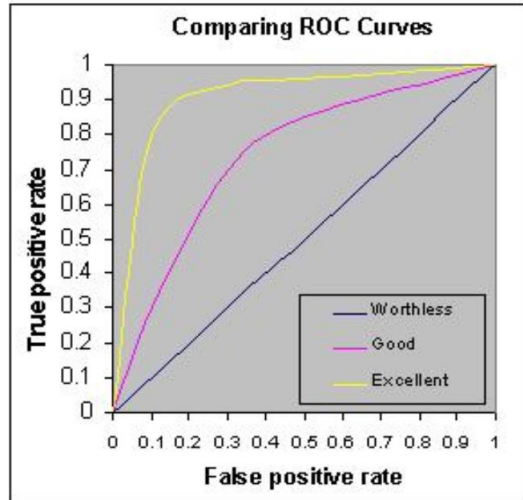


$$TPR = \frac{TP}{TP + FN} \quad FPR = \frac{FP}{TN + FP}$$

- Typically, a binary classifier has a decision boundary of  $\tau = 0.5$ 
  - If  $P(\text{sample has positive label}) \geq \tau$ , then predict a positive label for the sample. Else, predict a negative label.
- An ROC curve is traced by tracking (FPR, TPR) for a variety of classification thresholds  $\tau$  in the range  $[0, 1]$ 
  - All ROC curves pass through the points (1,1) and (0,0)
  - $\tau = 0 \Rightarrow$  Only positive labels predicted  $\Rightarrow$  (FPR, TPR) = (1,1)
  - $\tau = 1 \Rightarrow$  Only negative labels predicted  $\Rightarrow$  (FPR, TPR) = (0,0)

# ROC and AUC

Receiver Operating Characteristic curve: A plot of **False Positive Rate** Vs **True Positive Rate**



An ML model has parameters and for different values of the parameters, the model gives different FPR and TPR.

An **ROC (Receiver Operating Characteristic) curve** can be plotted for these different settings. A **good setting** of the parameters i.e. a good classifier, would have a **high TPR and low FPR**.

$$TPR = \frac{TP}{TP + FN}$$

$$FPR = \frac{FP}{TN + FP}$$

## Area Under Curve (AUC) in [0,1]:

A good classifier would have the largest area under the ROC curve.

The random predictor would have an ROC of 0.5 (the curve of the 'worthless' model in fig.)

## AUC value **interpretation** –

If a pair of data samples are drawn independently, one each from the positive and negative sets, AUC gives the probability that the classifier will predict a lower score for the negative sample as compared to the positive sample

# Be Careful with Interpreting Metrics...

- When comparing accuracies/AUCs between classifiers, make sure that they are operating on the same (or sufficiently similar) datasets
  - Otherwise, the implicit bias in the datasets can affect prediction outcomes
- If the AUC scores for two classifiers are similar, the individual class probabilities reported for samples by each classifier can be examined to determine which classifier is more discriminative
- AUC is only used with binary classifications
  - For multinomial classifications with AUC, a “one vs rest” scheme may be needed



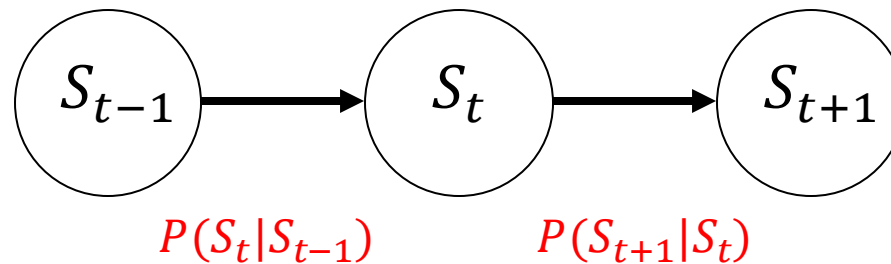
# Reinforcement Learning

Adapted from slides by Sergey Levine  
Dept of eecs UC Berkeley

---

# From Markov Chains to RL: Markov Model

We can use causal relationships to specify **conditional probabilities** (like in Bayesian networks)

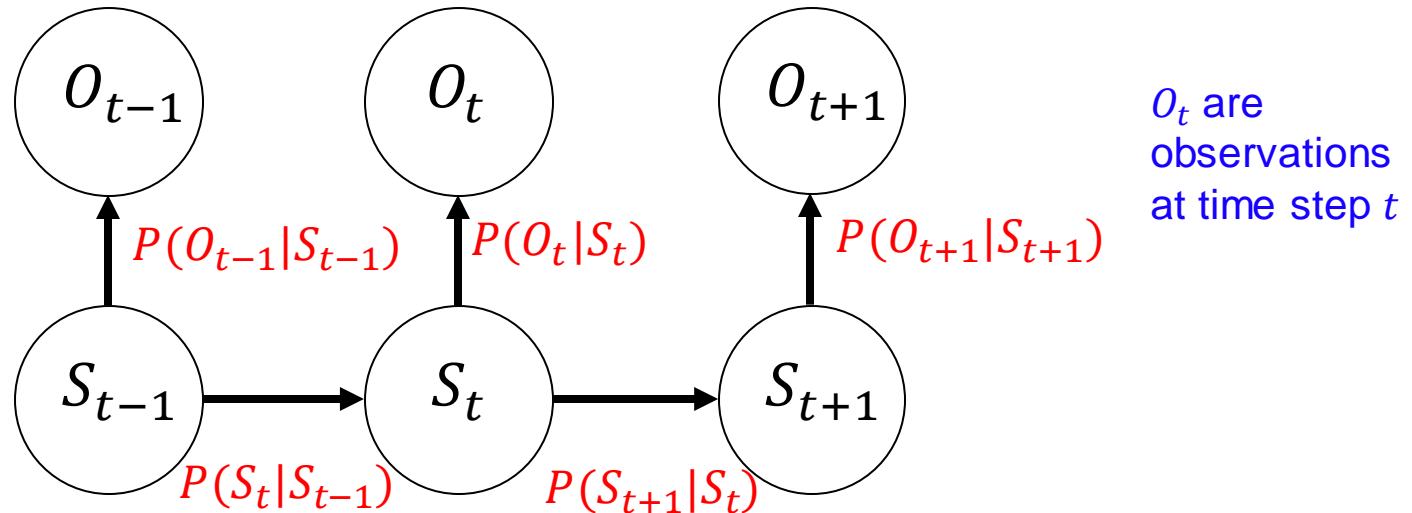


$S_t$  is the state  
at time step  $t$

Goal: Determine state using  $P(S_{t+1}|S_t)$

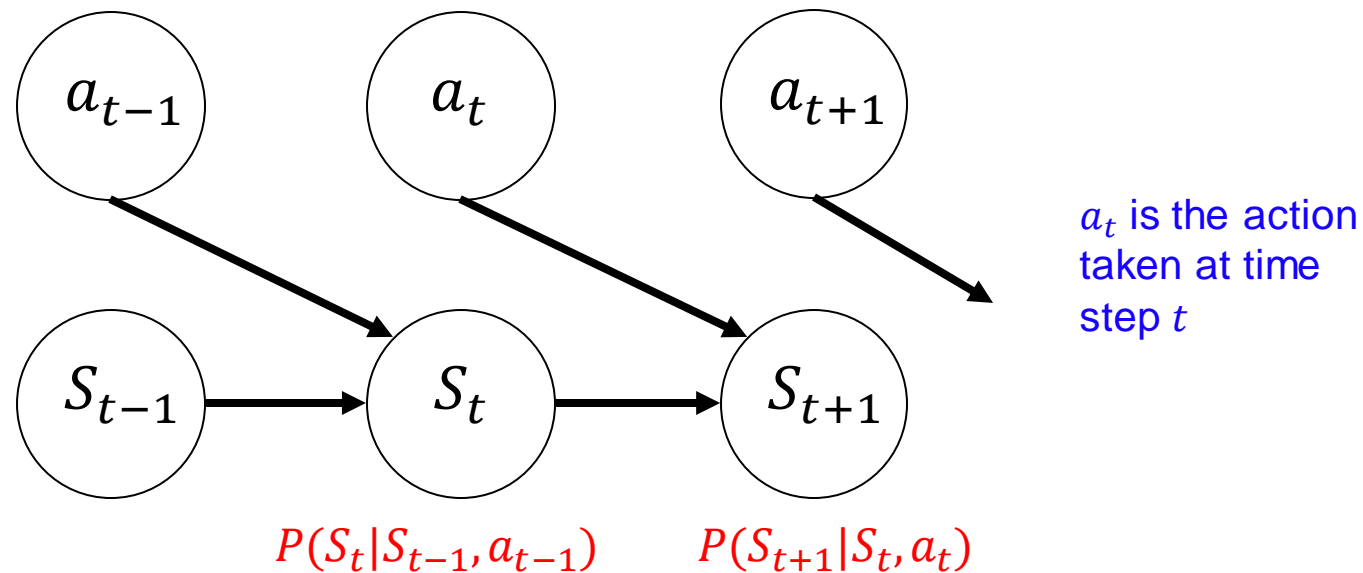


# From Markov Chains to RL: HMM



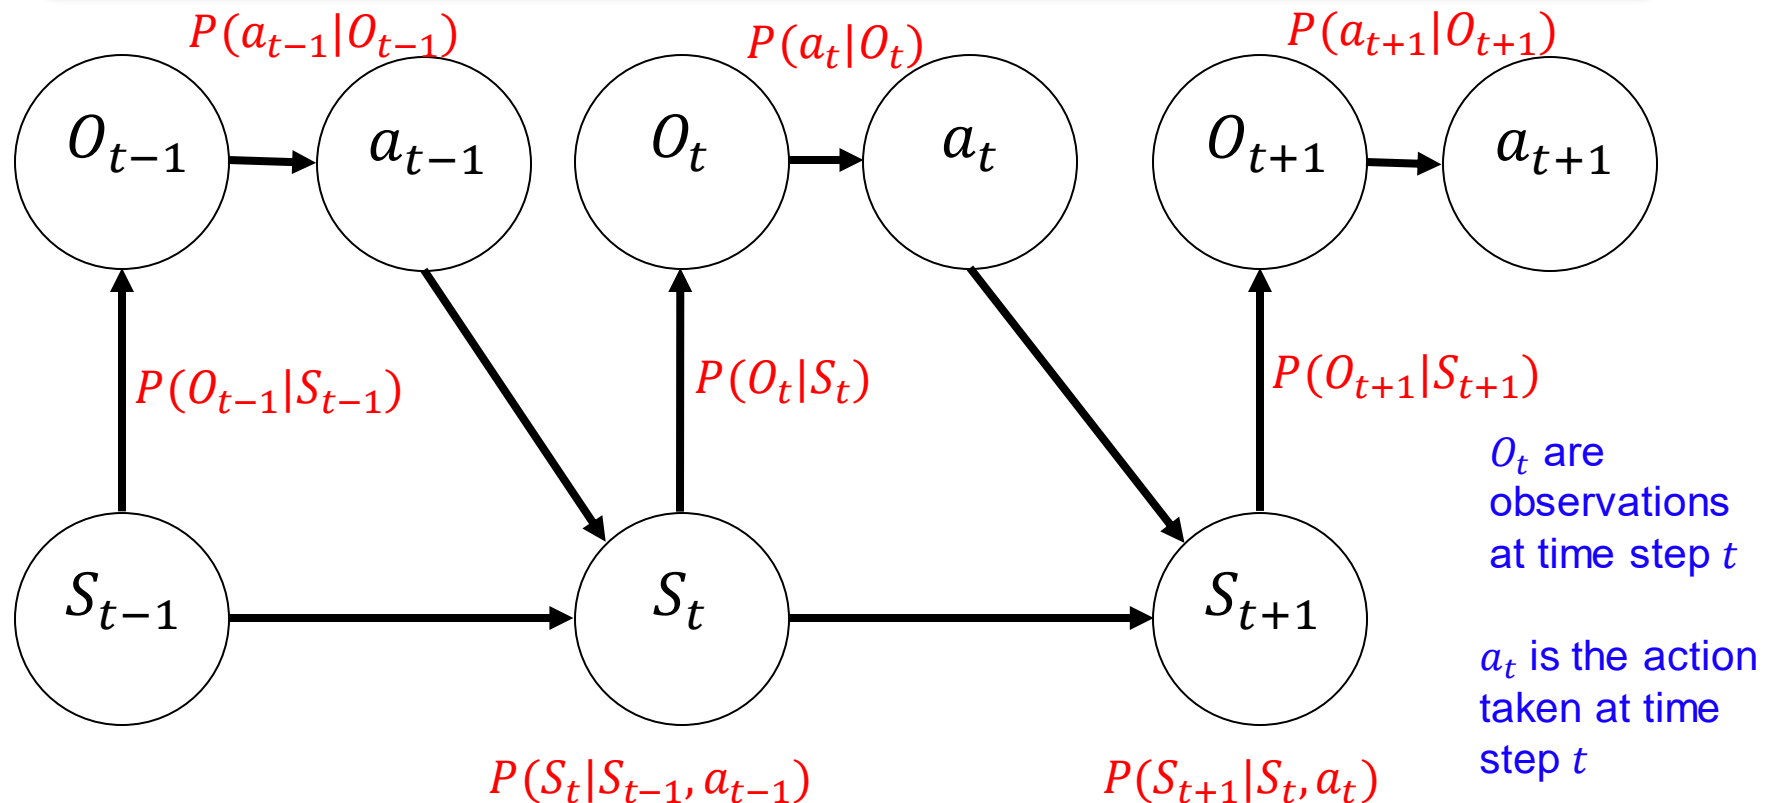
Goal: Determine hidden states using  $P(S_{t+1}|S_t), P(O_{t+1}|S_{t+1})$

# From Markov Chains to RL: Markov Decision Process



Goal: Identify a **policy**  $\pi_\theta(a_t|S_t)$ , which is the action  $a_t$  to take at state  $S_t$   
In an MDP, we **fully observe**  $S_t$

# From Markov Chains to RL: Partially Observed Markov Decision Process



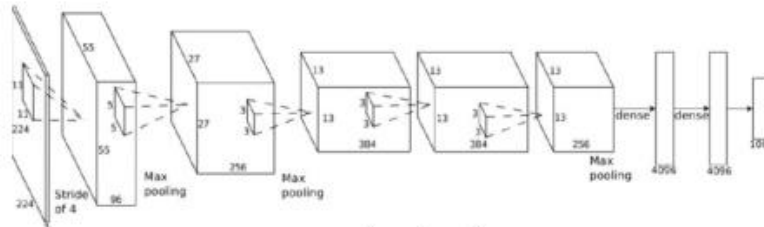
Goal: Identify a **policy**  $\pi_\theta(a_t|O_t)$ , which is the action  $a_t$  to take at state  $S_t$

In an POMDP, we **observe**  $O_t$  ( $S_t$  is hidden)

# Terminology & Notation



$\mathbf{o}_t$



$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$



$\mathbf{a}_t$

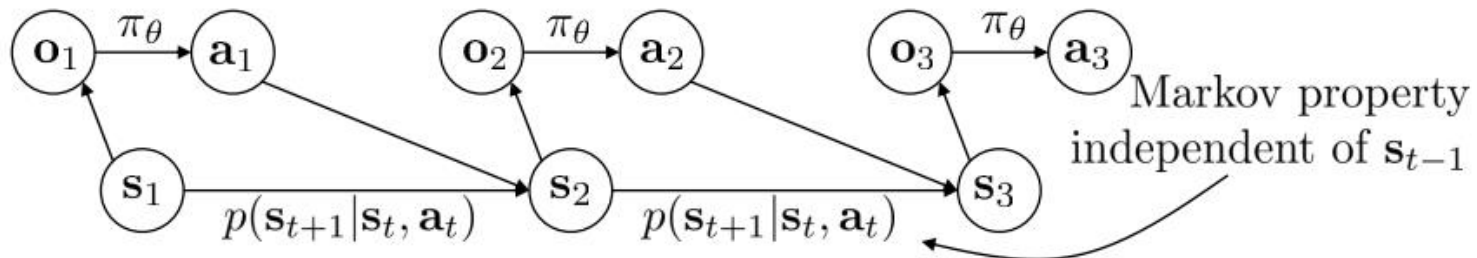
$\mathbf{s}_t$  – state

$\mathbf{o}_t$  – observation

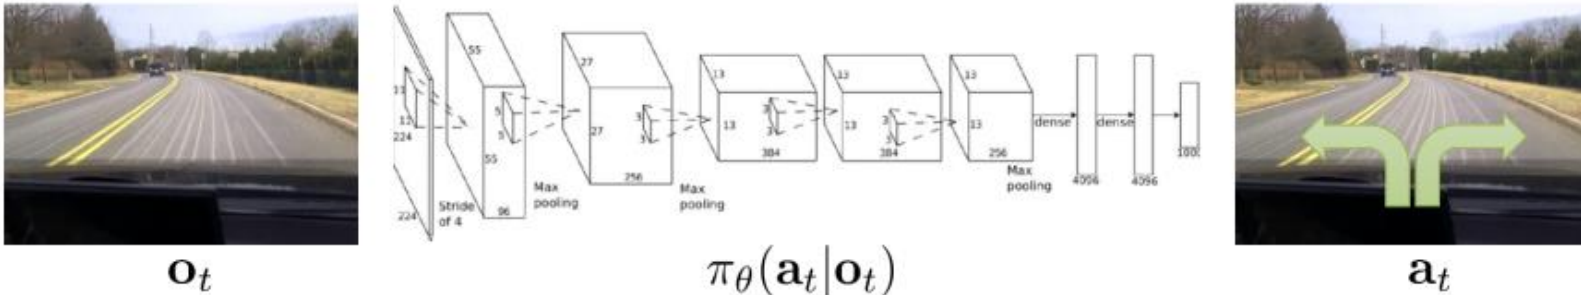
$\mathbf{a}_t$  – action

$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$  – policy

$\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$  – policy (fully observed)



# Reward functions



which action is better or worse?

$r(\mathbf{s}, \mathbf{a})$ : reward function

tells us which states and actions are better

$\mathbf{s}$ ,  $\mathbf{a}$ ,  $r(\mathbf{s}, \mathbf{a})$ , and  $p(\mathbf{s}' | \mathbf{s}, \mathbf{a})$  define

Markov decision process



high reward



low reward

# Definitions

Markov decision process

$$\mathcal{M} = \{\mathcal{S}, \mathcal{A}, \mathcal{T}, r\}$$

$\mathcal{S}$  – state space

states  $s \in \mathcal{S}$  (discrete or continuous)

$\mathcal{A}$  – action space

actions  $a \in \mathcal{A}$  (discrete or continuous)

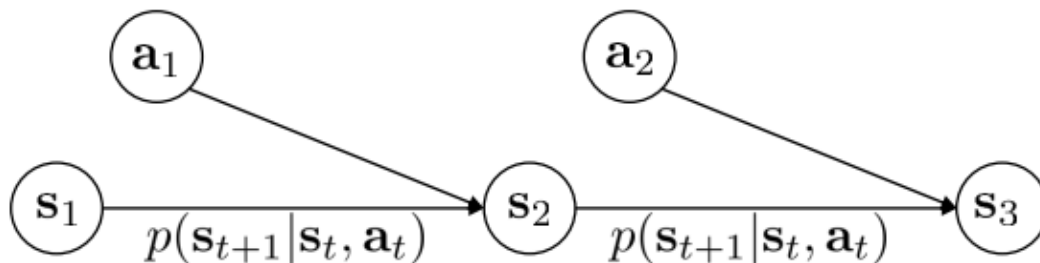
$\mathcal{T}$  – transition operator (now a tensor!)

let  $\mu_{t,j} = p(s_t = j)$

let  $\xi_{t,k} = p(a_t = k)$

let  $\mathcal{T}_{i,j,k} = p(s_{t+1} = i | s_t = j, a_t = k)$

$$\begin{aligned}\mu_{t+1,i} &= \sum_{j,k} \mathcal{T}_{i,j,k} \mu_{t,j} \xi_{t,k} \\ &= \sum_{j,k} p(s_{t+1} = i | s_t = j, a_t = k) p(s_t = j) p(a_t = k) \\ &= \sum_{j,k} p(s_{t+1} = i, s_t = j, a_t = k)\end{aligned}$$



All possible ways of  
getting  $s_{t+1} = i$

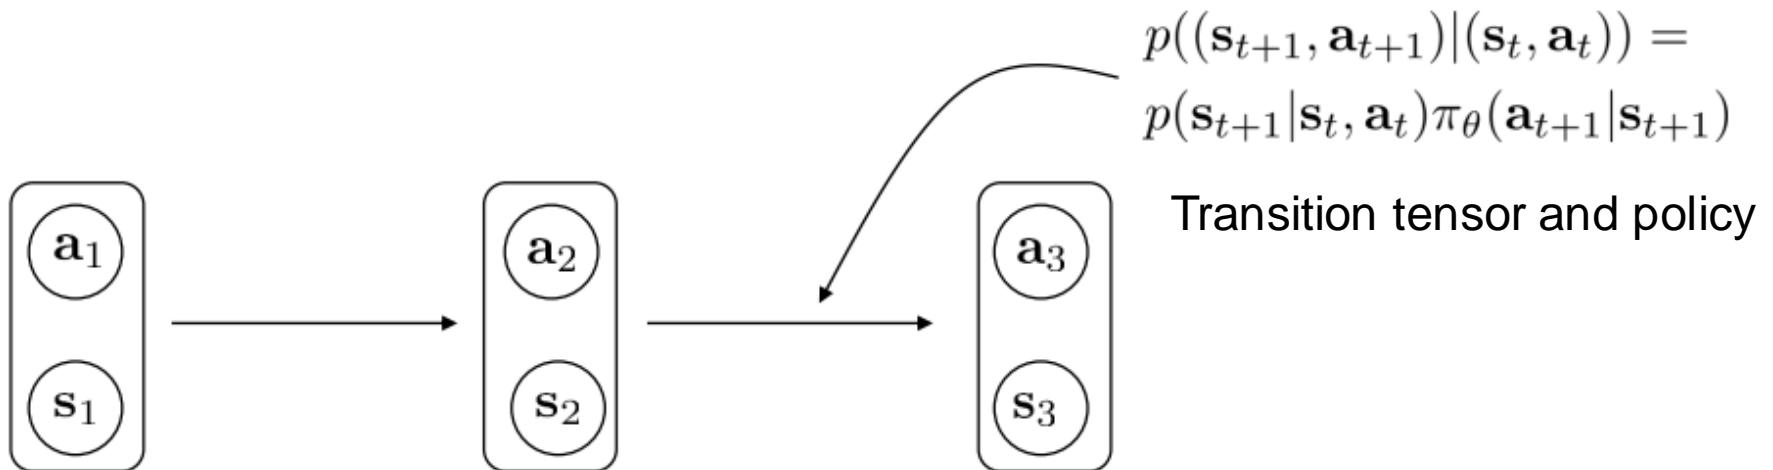
# Finite horizon case: state-action marginal

$$\theta^* = \arg \max_{\theta} E_{\tau \sim p_{\theta}(\tau)} \left[ \sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right]$$

Expected cumulative reward under probability distribution  $p_{\theta}(\mathcal{T})$

$$= \arg \max_{\theta} \sum_{t=1}^T E_{(\mathbf{s}_t, \mathbf{a}_t) \sim p_{\theta}(\mathbf{s}_t, \mathbf{a}_t)} [r(\mathbf{s}_t, \mathbf{a}_t)]$$

$p_{\theta}(\mathbf{s}_t, \mathbf{a}_t)$  state-action marginal



# Infinite horizon case: stationary distribution

$$\theta^* = \arg \max_{\theta} \sum_{t=1}^T E_{(\mathbf{s}_t, \mathbf{a}_t) \sim p_{\theta}(\mathbf{s}_t, \mathbf{a}_t)} [r(\mathbf{s}_t, \mathbf{a}_t)]$$

what if  $T = \infty$ ?

does  $p(\mathbf{s}_t, \mathbf{a}_t)$  converge to a *stationary* distribution?

$$\mu = \mathcal{T}\mu$$

stationary = the  
same before and  
after transition

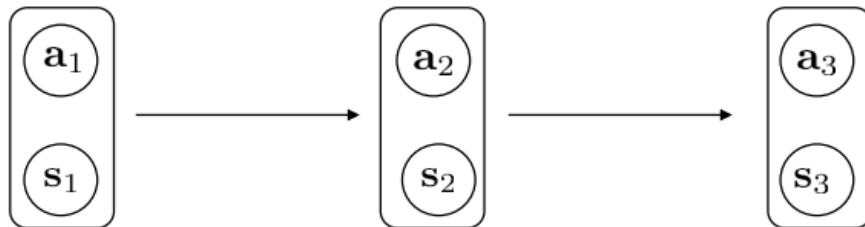
$$(\mathcal{T} - \mathbf{I})\mu = 0$$

$\mu$  is eigenvector of  $\mathcal{T}$  with eigenvalue 1!

(always exists under some regularity conditions)

$$\mu = p_{\theta}(\mathbf{s}, \mathbf{a}) \quad \text{stationary distribution}$$

state-action transition operator

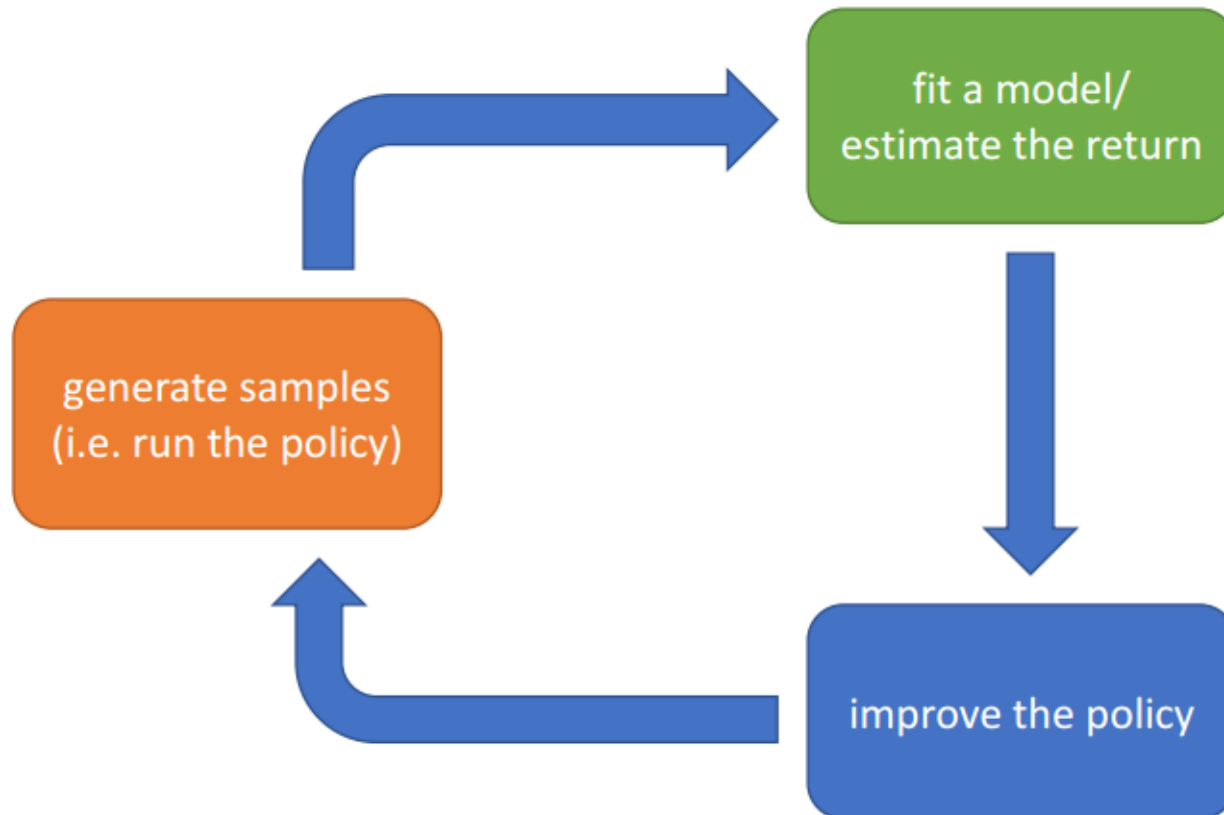


$$\begin{pmatrix} \mathbf{s}_{t+1} \\ \mathbf{a}_{t+1} \end{pmatrix} = \mathcal{T} \begin{pmatrix} \mathbf{s}_t \\ \mathbf{a}_t \end{pmatrix} \quad \begin{pmatrix} \mathbf{s}_{t+k} \\ \mathbf{a}_{t+k} \end{pmatrix} = \mathcal{T}^k \begin{pmatrix} \mathbf{s}_t \\ \mathbf{a}_t \end{pmatrix}$$

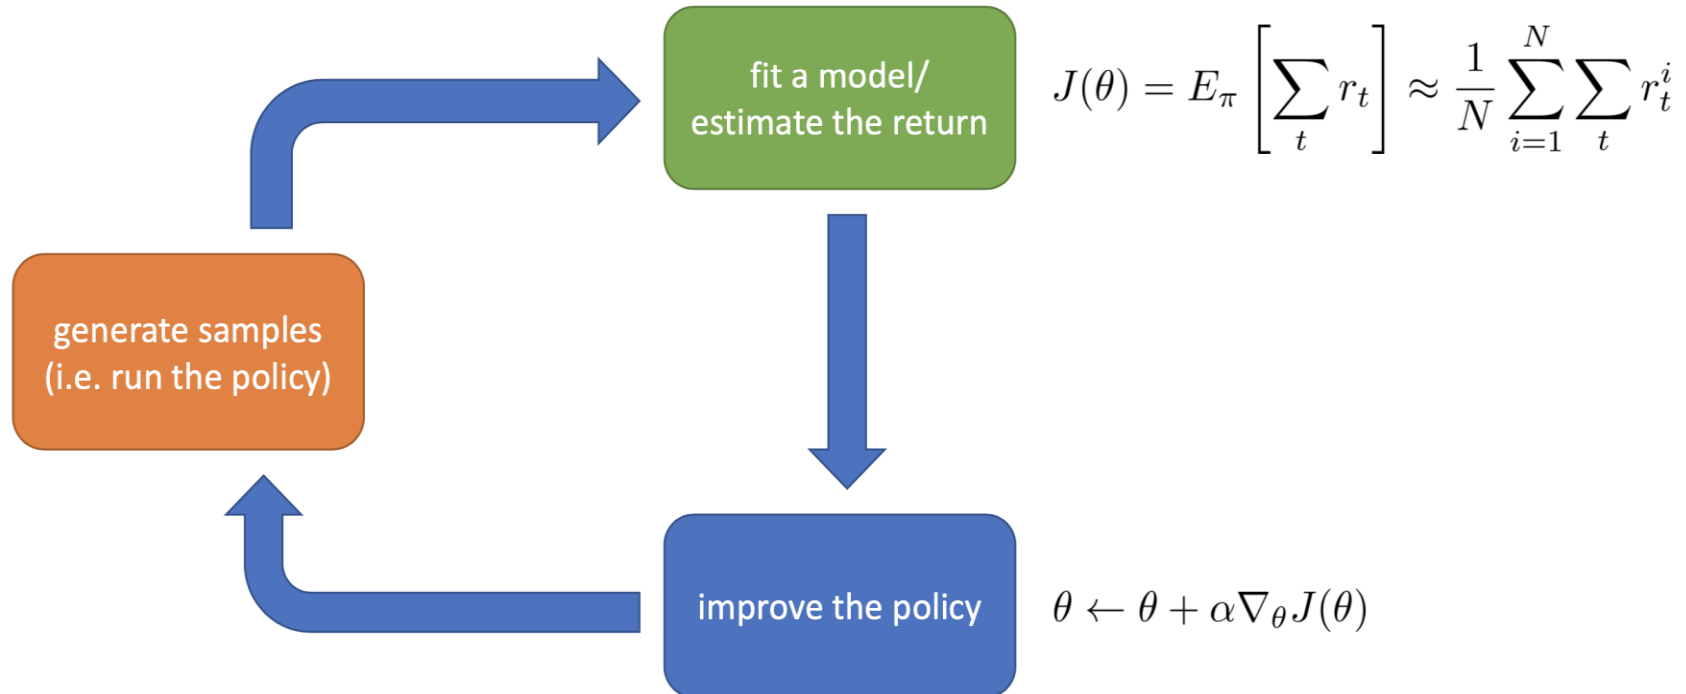
- The computation of the expected reward is difficult since  $p_{\theta}(\mathbf{s}_t, \mathbf{a}_t)$  changes at each  $t$
- By taking  $T \rightarrow \infty$ ,  $p_{\theta}(\mathbf{s}_t, \mathbf{a}_t) \rightarrow \mu = p_{\theta}(\mathbf{s}, \mathbf{a})$ ; now it is easier to compute the expected reward



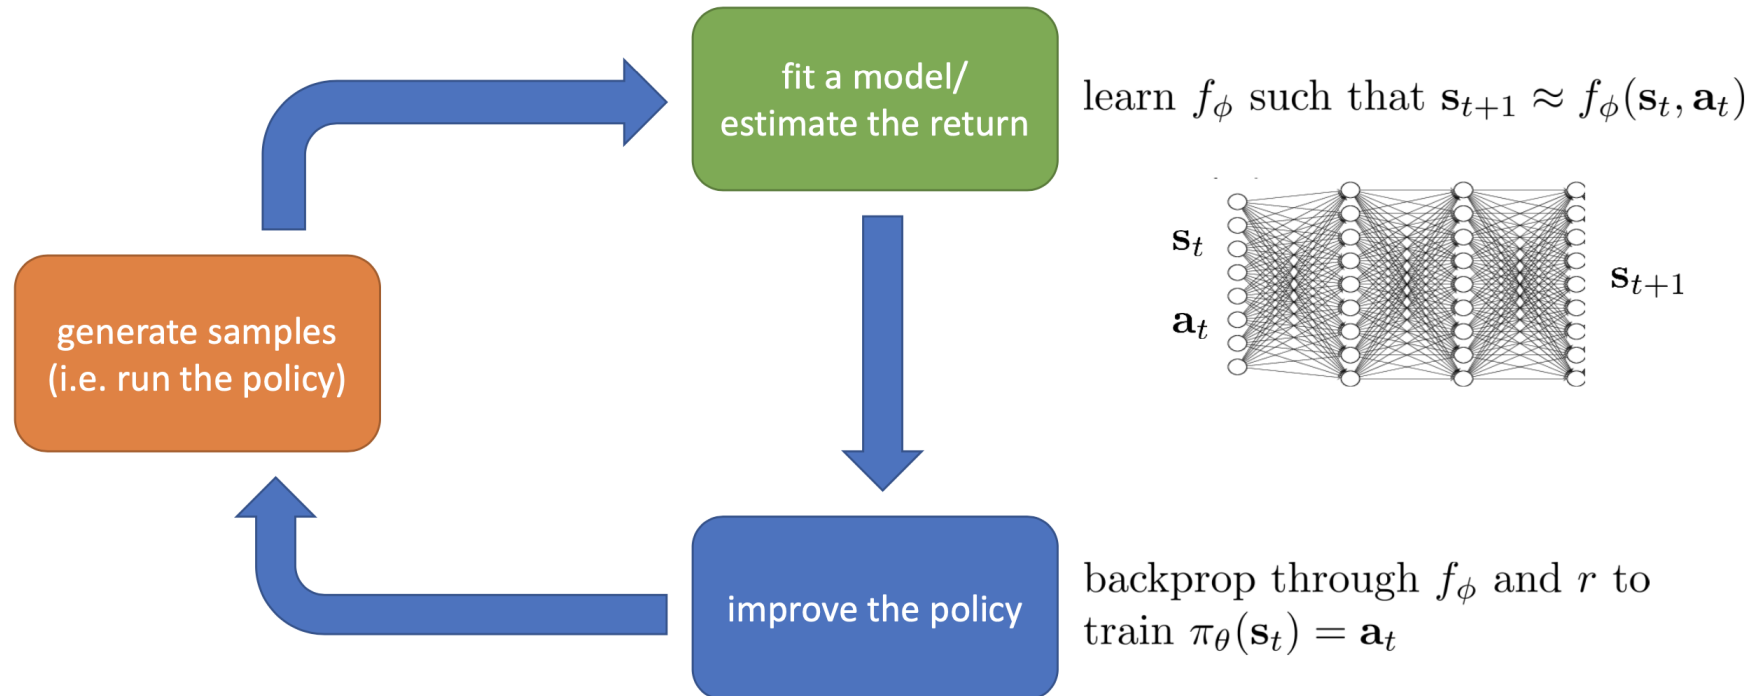
# The anatomy of a reinforcement learning algorithm



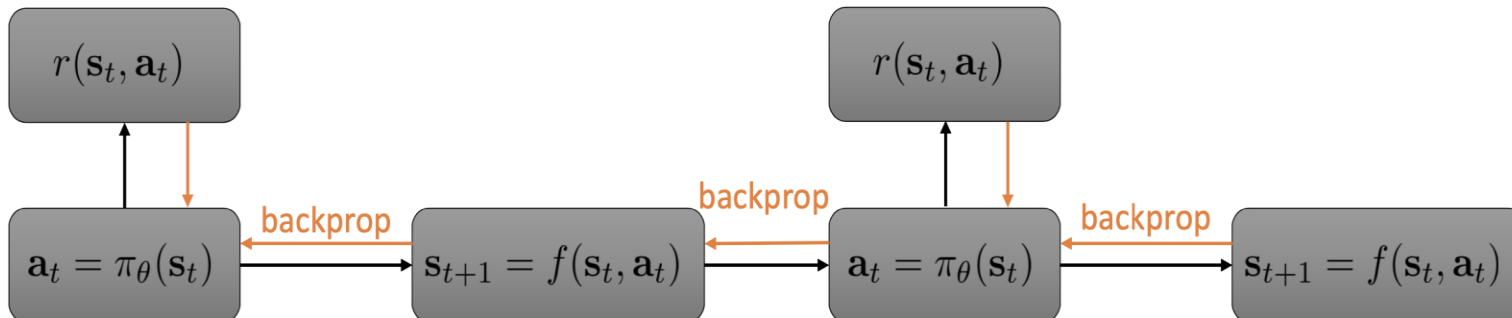
# A Simple RL Example



# Backpropagation in RL



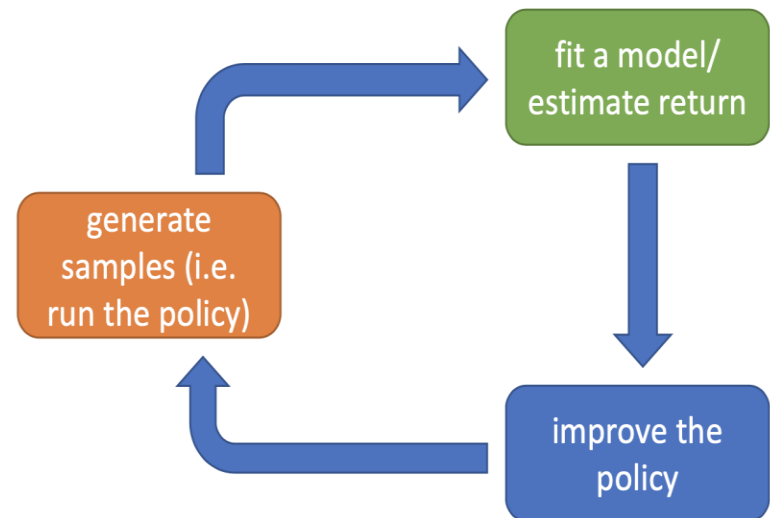
# Backpropagation in RL: Example



collect data

update the model  $f$

update the policy with backprop



# Types of RL algorithms

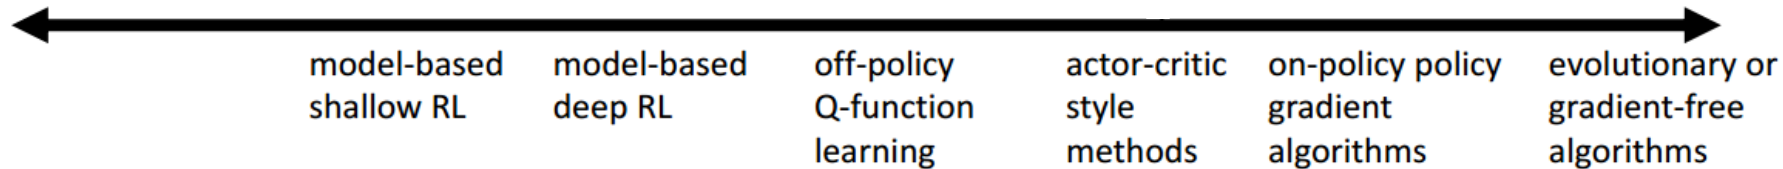
$$\theta^* = \arg \max_{\theta} E_{\tau \sim p_{\theta}(\tau)} \left[ \sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right]$$

- Policy gradients: directly differentiate the above objective
- Value-based: estimate value function or Q-function of the optimal policy (no explicit policy)
- Actor-critic: estimate value function or Q-function of the current policy, use it to improve policy
- Model-based RL: estimate the transition model, and then...
  - Use it for planning (no explicit policy)
  - Use it to improve a policy
  - Something else

# Comparison: sample efficiency

More efficient  
(fewer samples)

Less efficient  
(more samples)



Why would we use a *less* efficient algorithm?